

Department of CSE Artificial Intelligence

Experiment No. 2

Title: Study and Implement Multiple Linear Regression

Objectives:

1. Develop a machine learning model to predict calories burned based on exercise-related metrics.
2. Perform data preprocessing (handling missing values, encoding categorical variables).
3. Train and evaluate a Linear Regression model for calorie prediction.
4. Measure model performance using Mean Squared Error (MSE) and R^2 score.
5. Visualize actual vs. predicted values to assess model accuracy.

Problem Statement:

This practical builds a data-driven approach using linear regression to predict calorie expenditure based on exercise parameters such as duration, pulse rate, and max pulse.

Outcomes:

1. A trained Linear Regression model that predicts calorie expenditure.
2. Evaluation metrics (MSE, R^2) to assess model accuracy.
3. A scatter plot of actual vs. predicted values to visualize prediction quality.

Tools Required: 4GB RAM, Anaconda, Notebook

Theory:

Multiple Linear Regression (MLR) is an extension of simple linear regression, where a dependent variable (Y) is predicted based on multiple independent variables (X_1, X_2, \dots, X_n). The equation of MLR is:

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_n X_n + \epsilon$$

Where:

- Y = Dependent variable
- X_1, X_2, \dots, X_n = Independent variables
- β_0 = Intercept
- $\beta_1, \beta_2, \dots, \beta_n$ = Coefficients of independent variables
- ϵ = Error term

The model learns these coefficients by minimizing the Mean Squared Error (MSE):

$$MSE = \frac{1}{n} \sum_{i=1}^n (Y - \hat{Y})^2$$

where Y is the actual value and \hat{Y} is the predicted value.

The goodness-of-fit is measured by R^2 score:

$$R^2 = 1 - \frac{SS_{res}}{SS_{tot}}$$

where SS_{res} is the residual sum of squares and SS_{tot} is the total sum of squares.

Algorithm:

Step 1: Load Dataset

- Read data.csv using pandas.
- Display the first few rows to understand the structure.

Step 2: Data Preprocessing

- Handle missing values (drop or impute).
- Select independent (X) and dependent (y) variables.
- Convert categorical features using `pd.get_dummies()`.

Step 3: Split Data

- Use `train_test_split()` to split into training (80%) and testing (20%) sets.

Step 4: Train Linear Regression Model

- Initialize `LinearRegression()` from sklearn.
- Fit the model using `X_train` and `y_train`.

Step 5: Make Predictions

- Use `model.predict(X_test)` to generate predictions.

Step 6: Evaluate Model

- Compute MSE and R^2 score to measure accuracy.

Source Code:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
df = pd.read_csv("data.csv") # Replace with actual dataset
print(df.head())
df = df.dropna()
X = df[['Duration', 'Pulse', 'Maxpulse']] # Replace with actual feature names
y = df['Calories']
X=pd.get_dummies(X, drop_first=True)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
model = LinearRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f"Mean Squared Error: {mse}")
print(f"R-squared Value: {r2}")
plt.scatter(y_test, y_pred)
plt.xlabel("Actual Values")
```

```
plt.ylabel("Predicted Values")
plt.title("Actual vs Predicted Values")
plt.show()
```

Output:

Duration Pulse Maxpulse Calories

0 60 110 130 409.1

1 60 117 145 479.0


2 60 103 135 340.0

3 45 109 175 282.4

4 45 117 148 406.0

Mean Squared Error: 4354.413282161514

R-squared Value: 0.8664689420685995



Conclusion:
